

大数据与机器智能

哈希表与哈希函数

许书畅

2019年10月

Outline







- Hash tables are another sort of data structure that allows fast **INSERT/DELETE/SEARCH**.
 - like self-balancing binary trees
 - The difference is we can get better performance in expectation by using randomness.
- **Hash families** are the magic behind hash tables.
- **Universal hash families** are even more magical.

Goal:

Just like last time

- We are interesting in putting nodes with keys into a data structure that supports fast **INSERT/DELETE/SEARCH**.

- **INSERT** 
- **DELETE** 
- **SEARCH** 



Today:

- Hash tables:

- $O(1)$ expected time INSERT/DELETE/SEARCH

- Worse worst-case performance, but often great in practice.



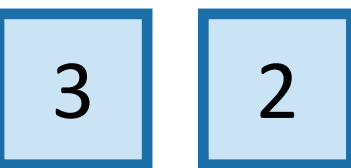
One way to get $O(1)$ time


This is called
“direct addressing”

- Say all keys are in the set $\{1, 2, 3, 4, 5, 6, 7, 8, 9\}$.

- **INSERT:** 

- **DELETE:** 

- **SEARCH:** 




That should look familiar

- Kind of like BUCKETSORT from Lecture 6.
- Same problem: if the keys may come from a “universe” $U = \{1, 2, \dots, 10000000000\} \dots$





*The universe is
really big!*



Solution?

Put things in buckets based on one digit

INSERT:




Now **SEARCH**



It's in this bucket somewhere...
go through until we find it.

Problem

INSERT:



Now **SEARCH**

22

....this hasn't made
our lives easier...




Hash tables

- That was an example of a hash table.
 - not a very good one, though.
- We will be **more clever** (and less deterministic) about our bucketing.
- This will result in fast (expected time) INSERT/DELETE/SEARCH.

But first! Terminology.



- We have a universe U , of size M .
 - M is really big.
- But only a few (say at most n for today's lecture) elements of M are ever going to show up.
 - M is waaaayyyyyyy bigger than n .
- But we don't know which ones will show up in advance.



A few elements are special and will actually show up.

Example: U is the set of all strings of at most 140 ascii characters. (128^{140} of them).

The only ones which I care about are those which appear as trending hashtags on twitter. [#hashinghashtags](#)

There are way fewer than 128^{140} of these.

Examples aside, I'm going to draw elements like I always do, as blue boxes with integers in them...

The previous example


with this terminology

- We have a universe U , of size M .
 - at most n of which will show up.
- M is waaaaayyyyyy bigger than n .
- We will put items of U into n buckets.
- There is a *hash function* $h: U \rightarrow \{1, \dots, n\}$ which says what element goes in what bucket.

For this lecture, I'm assuming that the number of things is the same as the number of buckets, both are n .

This doesn't have to be the case, although we do want:

#buckets = $O(\text{\#things which show up})$



This is a hash table (with chaining)


- Array of n buckets.
- Each bucket stores a linked list.
 - We can insert into a linked list in time $O(1)$
 - To find something in the linked list takes time $O(\text{length}(\text{list}))$.
- $h: U \rightarrow \{1, \dots, n\}$ can be any function:
 - but for concreteness let's stick with $h(x) = \text{least significant digit of } x$.

For demonstration purposes only!

This is a terrible hash function! Don't use this!




INSERT:



SEARCH 43:

Scan through all the elements in bucket $h(43) = 3$.




Sometimes this a **good idea**


Sometimes this is a **bad idea**

- How do we pick that function so that this is a good idea?

1. We want there to be not many buckets (say, n).
 - This means we don't use too much space
2. We want the items to be pretty spread-out in the buckets.
 - This means it will be fast to SEARCH/INSERT/DELETE



vs.



Worst-case analysis

- Goal: Design a function $h: U \rightarrow \{1, \dots, n\}$ so that:
 - No matter what input (fewer than n items of U) a bad guy chooses, the buckets will be balanced.
 - Here, balanced means $O(1)$ entries per bucket.
- If we had this, then we'd achieve our dream of $O(1)$
INSERT/DELETE/SEARCH

Can you come up with such a function?


This is impossible!



No deterministic hash function can defeat worst-case input!

We really can't beat the bad guy here.

- The universe U has M items
- They get hashed into n buckets
- At least one bucket has at least M/n items hashed to it.
- M is waayyyy bigger than n , so M/n is bigger than n .
- **Bad guy chooses n of the items that landed in this very full bucket.**



Solution: Randomness




The game




Plucky the pedantic penguin

What does **random** mean here? Uniformly random?


1. An adversary chooses any n items $u_1, u_2, \dots, u_n \in U$, and any sequence of INSERT/DELETE/SEARCH operations on those items.




INSERT 13, INSERT 22, INSERT 43,
INSERT 92, INSERT 7, SEARCH 43,
DELETE 92, SEARCH 7, INSERT 92




2. You, the algorithm, chooses a **random** hash function $h: U \rightarrow \{1, \dots, n\}$.



3. **HASH IT OUT** #hashpuns



Example



- Say that $h: U \rightarrow \{1, \dots, n\}$ is a uniformly random function.
 - That means that **$h(1)$** is a **uniformly random** number between 1 and n .
 - **$h(2)$** is also a **uniformly random** number between 1 and n , independent of $h(1)$.
 - **$h(3)$** is also a **uniformly random** number between 1 and n , independent of $h(1)$, $h(2)$.
 - ...
 - **$h(M)$** is also a **uniformly random** number between 1 and n , independent of $h(1)$, $h(2)$, ..., $h(M-1)$.

Randomness helps

Intuitively: The bad guy can't foil a hash function that he doesn't yet know.



Lucky the
Lackadaisical Lemur




Plucky the Pedantic
Penguin

Why not? What if there's some strategy that foils a random function with high probability?

We'll need to do some analysis...

What do we want?

It's **bad** if lots of items land in u_i 's bucket.
So we want **not that**.




More precisely

We could replace “2” here with any constant; it would still be good. But “2” will be convenient.

- We want:
 - For all ways a bad guy could choose u_1, u_2, \dots, u_n , to put into the hash table, and for all $i \in \{1, \dots, n\}$, $E[\text{number of items in } u_i \text{ 's bucket}] \leq 2$.
- If that were the case:
 - For each INSERT/DELETE/SEARCH operation involving u_i ,

$$E[\text{time of operation}] = O(1)$$

This is what we wanted at the beginning of lecture!

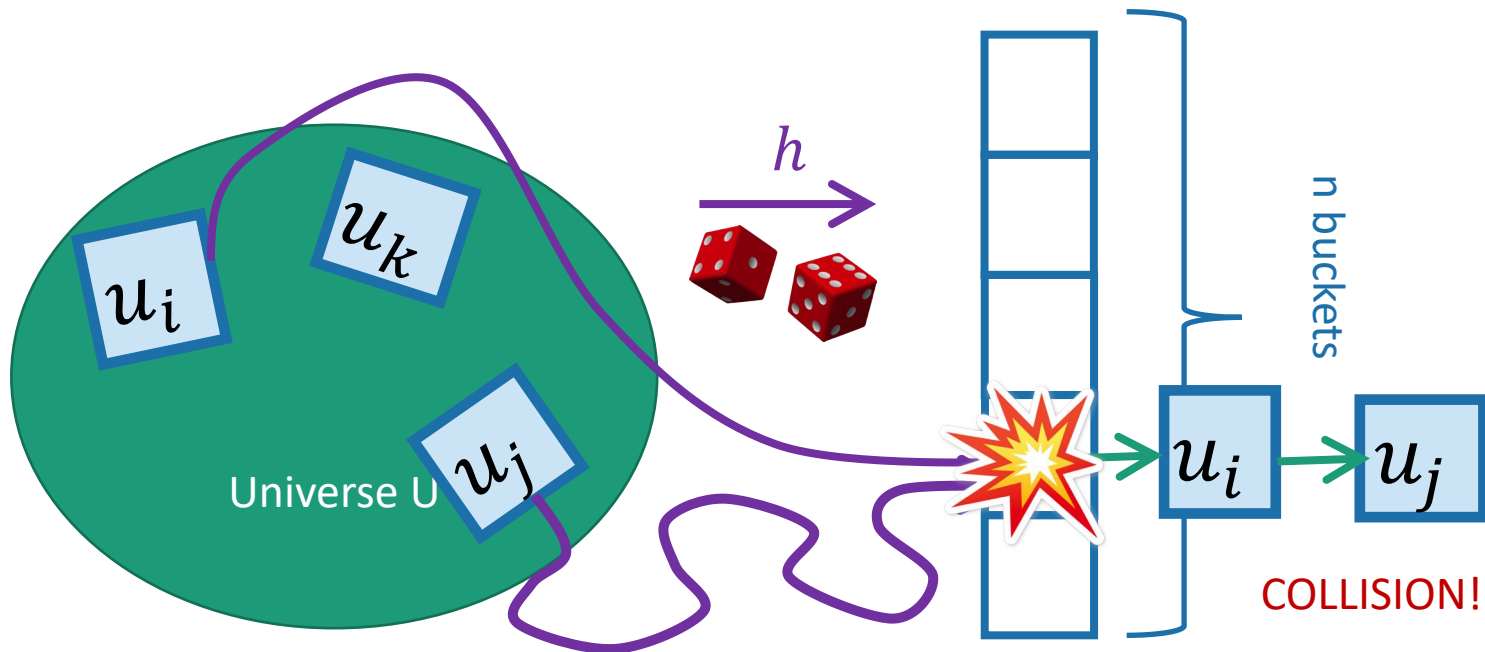


So we want:

- For all $i=1, \dots, n$,
 $E[\text{number of items in } u_i \text{'s bucket}] \leq 2.$

Expected number of items in u_i 's bucket?

- $E[\] = \sum_{j=1}^n P\{h(u_i) = h(u_j)\}$
- $= 1 + \sum_{j \neq i} P\{h(u_i) = h(u_j)\}$
- $= 1 + \sum_{j \neq i} 1/n$
- $= 1 + \frac{n-1}{n} \leq 2.$ That's what we wanted!




That's great!

- We just showed:
 - For all ways a bad guy could choose u_1, u_2, \dots, u_n , to put into the hash table, and for all $i \in \{1, \dots, n\}$,
$$E[\text{number of items in } u_i \text{'s bucket}] \leq 2.$$
- Which implies:
 - No matter what sequence of operations and items the bad guy chooses,
$$E[\text{time of INSERT/DELETE/SEARCH}] = O(1)$$
- So our solution is:

Pick a uniformly random hash function?

What's wrong with this plan?

- Hint: How would you implement (and store) and uniformly random function $h: U \rightarrow \{1, \dots, n\}$?



Think-Pair-Share Terrapins

- If h is a uniformly random function:
 - That means that $h(1)$ is a **uniformly random** number between 1 and n .
 - $h(2)$ is also a **uniformly random** number between 1 and n , independent of $h(1)$.
 - $h(3)$ is also a **uniformly random** number between 1 and n , independent of $h(1)$, $h(2)$.
 - ...
 - $h(M)$ is also a **uniformly random** number between 1 and n , independent of $h(1)$, $h(2)$, ..., $h(M-1)$.

A uniformly random hash function is not a good idea.

- In order to store/evaluate a uniformly random hash function, we'd use a lookup table:


All of the M things in the universe

| x | $h(x)$ |
|--------|--------|
| AAAAAA | 1 |
| AAAAAB | 5 |
| AAAAAC | 3 |
| AAAAAD | 3 |
| ... | |
| ZZZZZY | 7 |
| ZZZZZZ | 3 |

- Each value of $h(x)$ takes $\log(n)$ bits to store.
- Storing M such values requires $M\log(n)$ bits.
- In contrast, direct addressing (initializing a bucket for every item in the universe) requires only M bits....


Could we store a uniformly random h without using a lookup table?

- Maybe there's a different way to store h that uses less space?



Aside: description length

- Say I have a set S with s things in it.
- I get to write down the elements of S however I like, in binary using b bits.
- Then $b \geq \log(s)$:
 - There are 2^b binary strings of length b .
 - I need to have at least as many strings as I have items in S .
 - So $s \leq 2^b$ aka $b \geq \log(s)$




I'll call this one "Fido"
Or, 011011

This one is named "Hercules"
Or, 101111

We need $M\log(n)$ bits to store a random hash function $h:U \rightarrow \{1,\dots,n\}$

- Say that this elephant-shaped blob represents the set of all hash functions.
- It has size n^M . (Really big!)
- To write down a random hash function, we need $\log(n^M) = M\log(n)$ bits.
- A random hash function is just a random element in this set.




Technically we should argue that we need at least $M\log(n)$ bits on average when we draw a random hash function... why can't one item have a really long name and others have shorter names?


Solution

- Pick from a smaller set of functions.

A cleverly chosen subset of functions. We call such a subset a hash family.



We need only $\log |H|$ bits to store an element of H .




Outline

- **Hash tables** are another sort of data structure that allows fast **INSERT/DELETE/SEARCH**.
 - like self-balancing binary trees
 - The difference is we can get better performance in expectation by using randomness.
- **Hash families** are the magic behind hash tables.
- **Universal hash families** are even more magic.



Hash families

- A hash family is a collection of hash functions.





"All of the hash functions" is an example of a hash family.

Example:

a smaller hash family

This is still a terrible idea!
Don't use this example!
For pedagogical purposes only!

- $H = \{ \text{function which returns the least sig. digit,} \\ \text{function which returns the most sig. digit} \}$
- Pick h in H at random.
- Store just one bit to remember which we picked.



The game

h_0 = Most_significant_digit

h_1 = Least_significant_digit


$H = \{h_0, h_1\}$

1. An adversary (who knows H) chooses any n items $u_1, u_2, \dots, u_n \in U$, and any sequence of INSERT/DELETE/SEARCH operations on those items.


2. You, the algorithm, chooses a **random** hash function $h: U \rightarrow \{0, \dots, 9\}$. Choose it randomly from H .




I picked h_1



INSERT 19, INSERT 22, INSERT 42,
INSERT 92, INSERT 0, SEARCH 42,
DELETE 92, SEARCH 0, INSERT 92



3. HASH IT OUT #hashpuns



This is not a very good hash family

- $H = \{ \text{function which returns least sig. digit,} \\ \text{function which returns most sig. digit} \}$
- On the previous slide, the adversary could have been a lot more adversarial...

The game

h_0 = Most_significant_digit

h_1 = Least_significant_digit


$H = \{h_0, h_1\}$

1. An adversary (who knows H) chooses any n items $u_1, u_2, \dots, u_n \in U$, and any sequence of INSERT/DELETE/SEARCH operations on those items.




2. You, the algorithm, chooses a **random** hash function $h: U \rightarrow \{0, \dots, 9\}$. Choose it randomly from H .



I picked h_1



3. HASH IT OUT #hashpuns




Outline

- **Hash tables** are another sort of data structure that allows fast **INSERT/DELETE/SEARCH**.
 - like self-balancing binary trees
 - The difference is we can get better performance in expectation by using randomness.
- **Hash families** are the magic behind hash tables.
- **Universal hash families** are even more magic.



How to pick the hash family?


- Definitely not like in that example.
- Let's go back to that computation from earlier....



Expected number of items in u_i 's bucket?

- $E[\] = \sum_{j=1}^n P\{h(u_i) = h(u_j)\}$
- $= 1 + \sum_{j \neq i} P\{h(u_i) = h(u_j)\}$
- $= 1 + \sum_{j \neq i} 1/n$
- $= 1 + \frac{n-1}{n} \leq 2.$

All that we needed
was that this is $1/n$




Strategy

- Pick a small hash family H , so that when I choose h randomly from H ,

$$\text{for all } u_i, u_j \in U \text{ with } u_i \neq u_j, \\ P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$


In English: fix any two elements of U .
The probability that they collide under a random h in H is small.

- A hash family H that satisfies this is called a **universal hash family**.
- Then we still get $O(1)$ -sized buckets in expectation.
- But now the space we need is $\log(|H|)$ bits.
 - Hopefully pretty small!




So the whole scheme will be

Choose h randomly
from a **universal hash**
family H




We can store h in small space
since H is so small.



Universe U

h



Probably
these
buckets will
be pretty
balanced.

Universal hash family

- H is a ***universal hash family*** if, when h is chosen uniformly at random from H,

for all $u_i, u_j \in U$ with $u_i \neq u_j$,

$$P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

- Pick a small hash family H , so that when I choose h randomly from H ,

Example

$$\text{for all } u_i, u_j \in U \text{ with } u_i \neq u_j, \\ P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

- H is all of the functions $h: U \rightarrow \{1, \dots, n\}$
 - We saw this earlier – it corresponds to picking a uniformly random hash function.
 - Unfortunately this H is really really large.


- Pick a small hash family H , so that when I choose h randomly from H ,

Non-example

$$\text{for all } u_i, u_j \in U \text{ with } u_i \neq u_j, \\ P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

- $h_0 = \text{Most_significant_digit}$
- $h_1 = \text{Least_significant_digit}$
- $H = \{h_0, h_1\}$

Prove that this choice of H is
NOT a universal hash family!



- Pick a small hash family H , so that when I choose h randomly from H ,

Non-example

$$\text{for all } u_i, u_j \in U \text{ with } u_i \neq u_j, \\ P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

- h_0 = Most_significant_digit
- h_1 = Least_significant_digit
- $H = \{h_0, h_1\}$

NOT a universal hash family:

$$P_{h \in H} \{ h(101) = h(111) \} = 1 > \frac{1}{10}$$

A small universal hash family??

- Here's one:

- Pick a prime $p \geq M$.

- Define

$$f_{a,b}(x) = ax + b \mod p$$

$$h_{a,b}(x) = f_{a,b}(x) \mod n$$

- Define:

$$H = \{ h_{a,b}(x) : a \in \{1, \dots, p-1\}, b \in \{0, \dots, p-1\} \}$$

- Claim:


H is a universal hash family.



-
- The diagram illustrates the second step of a hash function: modular reduction. It shows two circular hash tables, U (green) and V (black), and a vertical array of three slots representing the final hash values.
- Step 1:** The green circle U contains elements 0, 1, 2, 3, 4. An arrow labeled $f_{2,1}(x)$ points from U to the black circle V . Below this arrow is the text: "This step just scrambles stuff up. No collisions here!".
- Step 2:** The black circle V contains elements 0, 1, 2, 3, 4. Arrows from V point to a vertical array of three slots. The arrows are labeled with the function $f_{2,1}(x)$ applied to the elements: $f_{2,1}(0)$, $f_{2,1}(1)$, $f_{2,1}(2)$, $f_{2,1}(3)$, and $f_{2,1}(4)$. Below this array is the text: "This step is the one where two different elements might collide.".
- Modular Reduction:** An arrow labeled "mod 3" points from the array of slots to the final result. A starburst icon indicates a collision where two different elements map to the same slot.

Ignoring why this is a good idea

- Can we store h with small space?




- Just need to store two numbers:
 - a is in $\{1, \dots, p-1\}$
 - b is in $\{0, \dots, p-1\}$
 - So about $2\log(p)$ bits
 - By our choice of p , that's $O(\log(M))$ bits.

Compare: direct addressing was M bits!

Twitter example: $\log(M) = 140 \log(128) = \mathbf{980}$ vs $M = \mathbf{128^{140}}$


But let's check that it does work

M=200, n=10




So the whole scheme will be

Choose a and b at random
and form the function $h_{a,b}$




We can store h in space
 $O(\log(M))$ since we just need
to store a and b .



Universe U

$h_{a,b}$



Probably
these
buckets will
be pretty
balanced.

Outline

- **Hash tables** are another sort of data structure that allows fast **INSERT/DELETE/SEARCH**.
 - like self-balancing binary trees
 - The difference is we can get better performance in expectation by using randomness.
- **Hash families** are the magic behind hash tables.
- **Universal hash families** are even more magic.

Recap 

Want $O(1)$

INSERT/DELETE/SEARCH

- We are interesting in putting nodes with keys into a data structure that supports fast INSERT/DELETE/SEARCH.

- INSERT


5

- DELETE

4

- SEARCH

52

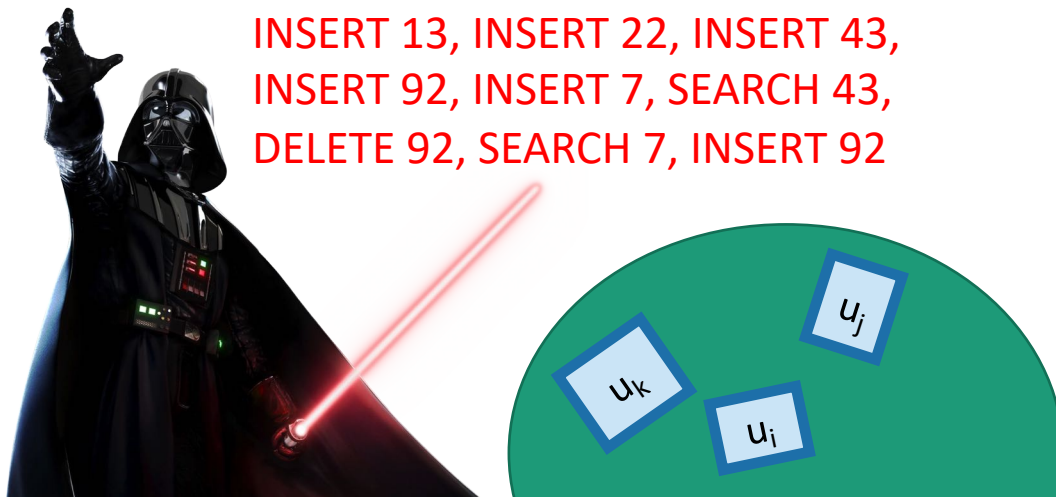


We studied this game


1. An adversary chooses any n items $u_1, u_2, \dots, u_n \in U$, and any sequence of L INSERT/DELETE/SEARCH operations on those items.




INSERT 13, INSERT 22, INSERT 43,
INSERT 92, INSERT 7, SEARCH 43,
DELETE 92, SEARCH 7, INSERT 92



2. You, the algorithm, chooses a **random** hash function $h: U \rightarrow \{1, \dots, n\}$.



3. HASH IT OUT



Uniformly random h was good

- If we choose h uniformly at random,


for all $u_i, u_j \in U$ with $u_i \neq u_j$,

$$P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

- That was enough to ensure that all INSERT/DELETE/SEARCH operations took $O(1)$ time in expectation, even on adversarial inputs.

Uniformly random h was bad

- If we actually want to implement this, we have to store the hash function h .
- That takes a lot of space!
 - We may as well have just initialized a bucket for every single item in U .
- Instead, we chose a function randomly from a smaller set.



We needed a **smaller set**
that still has this property

- If we choose h uniformly at random in H ,
for all $u_i, u_j \in U$ with $u_i \neq u_j$,
$$P_{h \in H} \{ h(u_i) = h(u_j) \} \leq \frac{1}{n}$$

This was all we needed to make
sure that the buckets were
balanced in expectation!

- We call any set with that property a
universal hash family.
- We gave an example of a really small one 😊



Hashing a universe of size M into n buckets, where at most n of the items in M ever show up.

Conclusion:

- We can build a hash table that supports **INSERT/DELETE/SEARCH** in $O(1)$ expected time
- Requires $O(n \log(M))$ bits of space.
 - $O(n)$ buckets
 - $O(n)$ items with $\log(M)$ bits per item
 - $O(\log(M))$ to store the hash function